Adoption Agent Record

Xinnan Li | Individual assignment | November 3, 2026

A Agent and Version

Student / team: Xinnan Li / Not applicable - individual assignment

Assignment / specialist: Assignment 6 — Emerging Technology Adoption Agent

Date: November 3, 2026

Agent version: 0.2-student

GitHub repository: https://github.com/xl5921/masy1800-assignment6-adoption-agent

Branch: assignment-6-adoption

Commit: 7984f1f (agent implementation; short ID observed in user screenshot; full hash not independently retrieved)

Frozen scaffold: MASY1800_ET_Agent_Scaffold_v1_0

B Specialist Purpose and Context

Governing question: Should this organization adopt, experiment with, monitor, defer, or reject this technology for this application now?

Technology: LLM-based customer-service drafting with internal knowledge retrieval.

Application: Draft order, returns-policy and after-sales replies, reviewed and sent by staff. No autonomous financial actions or policy exceptions; order status requires authorized records.

Organization / industry: Two hypothetical mid-sized consumer-goods e-commerce retailers.

Posture: A is moderately proactive; B has lower risk tolerance, stricter sequential approvals and weaker knowledge maintenance.

Consequences: Incorrect commitments, invented order status, privacy disclosure and loss of trust. Human review does not eliminate these risks.

Decision / horizon: Whether to start a limited pilot within three months, not full deployment.

C ThreeLevel Finding

ET generally: Evidence supports potential benefits in assisted customer service, with heterogeneous effects and persistent accuracy, privacy and human-reliance risks. Broad individual adoption does not establish enterprise readiness. No specific current model has been assessed.

Application: Grounded drafting warrants evaluation for bounded tasks, but reliable policies, authorized order information, review quality and comparisons against simpler alternatives are essential.

Organization: A can prepare a verified offline PoC using its existing knowledge ownership; B first needs accountable ownership and policy reconciliation. Both lack evidence permitting live deployment now.

Bounded implication: A — PROOF_OF_CONCEPT after minimum reference/data checks. B — DEFER live trial and remediate, then PoC. A future live experiment requires technical gates and case-specific approvals.

Confidence / uncertainty: Moderate confidence in staged action, low confidence in local net benefits and production readiness. No company performance data exist.

Change conditions: Verified policies, satisfactory quality and handling-time comparisons, reviewed costs/data controls, and approvals could justify a bounded live pilot. Failures support remediation, simpler alternatives or continued deferral.

D Evidence and Testing

Sources: E1, published Generative AI at Work (2025); E2, NIST AI 600-1 (2024); E3, Microsoft RAG guidance (updated 2026-08-21); E4, St. Louis Fed adoption report (2025-11-13). URLs, passage locators and transfer limitations: shared_evidence.json. Both cases received identical evidence content.

Primary test: v1 recommended conditional BOUNDED_EXPERIMENT but required offline prerequisites. v2 makes the immediate category PROOF_OF_CONCEPT and explicitly denies present live authorization.

Contrast stable finding: General technology evidence and its limitations stayed the same.

Contrast changed finding: B's unresolved knowledge ownership and approvals justify remediation and deferral of live use. Multiple context factors changed, so differences cannot be attributed solely to risk appetite.

Observed weakness: v1's primary label obscured the immediate action; both outputs lacked operational pass thresholds and sampling detail.

Revision: Preserved v1 instructions/metadata, outputs and record. Added NOW / NOT AUTHORIZED NOW / CONDITIONAL NEXT ACTION, a proposed 120-case stratified offline design, comparator workflows, measurable quality and time gates, stop rules and a capped live-stage proposal.

Retest: Both v2 outputs passed the supplied basic validator. No extra top-level output fields were found. Frozen-core check passed. The response wording now distinguishes immediate and conditional actions and supplies measurable proposals.

Remaining limitations: Thresholds are unvalidated management proposals; small samples do not certify rare-event safety. Same-conversation generation is not independent replication. The 120/100-case business evaluations were proposed, not performed.

Abstention: No unconditional live approval, compliance assurance, positive ROI claim or full-deployment recommendation without missing local evidence.

Files: responses/primary_response_v1.json, contrast_1_response_v1.json, primary_response_v2.json, contrast_1_response_v2.json; records/first_test_record.json and second_test_record.json. Prompt packets are in the scaffold work folder.

E Independent Judgment and Handoff

Credibility for team consideration: Suitable as a provisional decision-support candidate because it separates evidence from assumptions, keeps general findings stable and identifies conditional organizational actions. Not a validated deployment authority.

Team recommendation: Retain staged evaluation, compare simpler alternatives, and require specific evidence before live exposure. Reproduce runs in separate sessions if independent comparison is important.

Main uncertainty: Local quality and net handling-effort benefit under actual knowledge and review conditions.

Reversal condition: Approved local testing and resolved governance prerequisites can move PoC/deferral toward a bounded pilot; serious failures move it back toward remediation or rejection for the present.

Handoff to Organizational Adoption Agent: Detailed training, workload/staffing, stakeholder engagement and rollout design. Legal/privacy approval remains with authorized reviewers.

AI / verification note: See the accompanying AI verification draft. Original sources were checked by the assistant; personal student source verification should not be claimed until completed.

Submission note: Repository metadata was completed locally after implementation commit 7984f1f. These metadata updates and accompanying submission reports are not claimed to be included in that GitHub commit. Baseline short commit: 31f57e1, observed in the earlier screenshot.